




Figure 2: Breakdown of GPT-4 judgment over contextual factors, as we progress through tiers 1, 2.a and 2.b.

Table 3: Overview of metric values for Tier 3. Lower is better for all metrics.

	Metric	GPT-4	ChatGPT	InstructGPT	Mixtral	Llama-2 Chat	Llama-2
Leak.	Leakage thru. String Match	0.22	0.93	0.79	0.96	1.00	0.99
	Leakage thru. Proxy Agent	0.20	0.89	0.74	0.83	0.99	0.96
ToM.	Information Access. Err.	0.04	0.40	0.76	0.98	1.00	1.00
	Private Info. Access. Err.	0.03	0.32	0.70	0.97	1.00	1.00
	Binary Control Question	0.06	0.04	0.00	0.04	0.99	1.00




Figure 3: Breakdown of the string matching leakage reported for GPT-4 in Tier 3, with respect to different contextual factors. Lower means lower leakage.

that for all the evaluations, we instruct the model to respond **while considering privacy norms** (§3.3). The average case metric results and results without any privacy-preserving instructions are in Appendix B.3.

Overall, we find the leakage is alarmingly high for the open source models, and even for ChatGPT. Furthermore, we observe the error rates are high for ToM and control questions (§ 3.3) in most models except GPT-4 and ChatGPT. This suggests that most models struggle to discern who has access to which information. The performance of ChatGPT and GPT-4 for those question types may indicate that they have some ability to follow the flow of information. However, the high leakage of the secret in the free-form response reveals their limitation in reasoning over such knowledge and adequately controlling the flow.

To further analyze the leakage, we plot a detailed breakdown of the results, for the best performing model, GPT-4 in Figure 3. We find information regarding sexual orientation/self-harm is most/least likely to be revealed on average, and the incentives of helping others/gaining money lead to most/least leakage. Also, we observe **contention** between different contextual factors. For instance, although the model tends to not reveal information about self-harm, it opts to do so when the incentive is helping others. We attribute this behaviour to the alignment of such models to be ‘helpful’ Bai et al. (2022). We provide detailed heatmaps for other models, with and without privacy-inducing prompts (i.e., without telling the model to adhere to privacy norms, which make it leak even more) in Appendix B.5.1.

Table 4: Overview of metric values for Tier 4, where models are used as AI meeting assistants generating meeting summary and personal action items. Lower is better for all metrics.

	Metric	GPT-4	ChatGPT	InstructGPT	Mixtral	Llama2 Chat	Llama 2
Act. Item	Leaks Secret (Worst Case)	0.80	0.85	0.75	0.85	0.90	0.75
	Leaks Secret	0.29	0.38	0.28	0.54	0.43	0.21
	Omits Public Information	0.76	0.89	0.84	0.93	0.86	0.93
	Leaks Secret or Omits Info.	0.89	0.96	0.91	0.98	0.95	0.96
Summary	Leaks Secret (Worst Case)	0.80	0.85	0.55	0.70	0.85	0.75
	Leaks Secret	0.39	0.57	0.09	0.28	0.35	0.21
	Omits Public Information	0.10	0.27	0.64	0.42	0.73	0.77
	Leaks Secret or Omits Info.	0.42	0.74	0.68	0.65	0.92	0.87

Table 5: Overview of metric values for Tiers 3 & 4, where the model is instructed to do chain of thought reasoning, as a possible mitigation. Lower is better for all metrics.

		w/o CoT		w/ CoT	
	Metric	GPT-4	ChatGPT	GPT-4	ChatGPT
Tier3	Leakage thru. String Match	0.22	0.93	0.24	0.95
	Leaks Secret	0.39	0.57	0.40	0.61
Tier4	Omits Public Information	0.10	0.27	0.21	0.39
	Leaks Secret or Omits Info.	0.42	0.74	0.52	0.83

4.4 TIER 4 RESULTS

Table 4 summarizes the results for Tier 4. We provide both average and worst-case results for the leakage metric, across 10 runs. We find that all models have relatively high levels of leakage, as they tend to regurgitate the secret discussed at the beginning of the meeting. This leakage is higher in the meeting summary task compared to the personal action item generation task. We hypothesize that this could be due to the model being instructed to generate the action item specifically for person X (who is not supposed to know the secret), whereas in the summary generation the model is instructed to generate summary for all attendees, hence it isn’t able to reason that X is among the attendees and that the secret should be withheld. Additionally, we also report an aggregated metric, where we consider the response erroneous if it either leaks the secret or misses an important public action item. We observe a high error rate across all models, including GPT-4. Break-down of the findings, manual inspection and error analysis and detailed results without the privacy instructions can be found in Appendix B.4. Additionally, to account for the possible ‘familiarity’ of GPT-4 with scenarios generated by itself, we generate the Tier 4 scenarios with ChatGPT as well, and evaluate GPT-4 and ChatGPT on it. Those results are also provided in the Appendix.

4.5 IS CHAIN OF THOUGHT REASONING A VIABLE MITIGATION?

In this section, we present the results for Tiers 3 & 4, but we prompt the model with chain of thought reasoning (Wei et al., 2022) as a possible mitigation. Specifically, we add the sequence ‘Take a deep breath² and work on this step by step.’ to the instruction provided to the model, as proposed in Yang et al. (2023). We keep the prior instructions to preserve privacy as well. We only consider the final response for our evaluations and do not take the reasoning steps into consideration. Table 5 shows the results of this experiment. We can see that for almost all tasks, using chain of thought (CoT) does not improve leakage, in fact it makes the leakage more severe.

Inspection of the generated chains and categorization of different error types. To better understand the reasons behind the failure of the LLMs, we manually inspected the chains of thought generated by ChatGPT for Tier 3. We found that the scenarios where models reveal the secret falls into the the following four categories, shown in Table 6. We can see that the majority of failures fall under the second category, where the model acknowledges the importance of privacy, but is not capable of applying/operationalizing this knowledge. The second most common failure is related to theory of mind, as the model assumes that Z has access to the secret despite the scenario explicitly stating that only X and Y shared the secret.

² Yang et al. (2023) find this prompt to be most effective, through a prompt optimization method.

Table 6: Categorization of different error types in Tier 3 chain of thoughts.

Category	Description	Redacted example	Rate(%)
1	No acknowledgment of privacy and just reveals X’s secret to Z	...I actually have some personal experience supporting someone who went through a similar struggle. Laura, our classmate, had an eating disorder because of her body dissatisfaction...	7.5
2	Does acknowledge privacy, but reveals the X’s secret to Z	...but I think it’s important to consider Jane’s privacy and the trust she placed in me by confiding in me about her affair...	51.2
3	Does acknowledge privacy, but reveals X’s secret while reassuring Y that this interaction between Y and Z will be a secret	Let me share with you a personal experience that James went through. He once took credit for someone else’s coding, [...] Your privacy is of utmost importance, and I will keep our conversation confidential.	2.5
4	Theory of mind failure (i.e., assumes Z knows information about the secret) and reveals the secret	As you know, Adam and I have maintained confidentiality about his transformation and recovery, and it’s essential that we continue to respect his privacy.	38.3

5 CONCLUSION AND DISCUSSION

We introduce CONFAIDE to investigate the risk of *contextual privacy* leaks in LLMs. We identify new shortcomings in terms of privacy reasoning and theory of mind, demonstrating that even models that have undergone intensive RLHF training still lack reasoning on what should and should not be shared in various contexts. Finally, we explore possible mitigations, showing that straightforward measures, such as fortifying the prompt by instructing the model to maintain privacy or using chain of thought reasoning, are insufficient. Altogether our results highlight that more fundamental solutions are needed for LLMs to safely preserve privacy while deployed in real-world applications.

Inference-time privacy definitions: Our findings also point to an existing gap in the literature, regarding privacy definitions at inference time, which can have serious consequences. For instance a recent prompt-injection attack reverse-engineered Bing Chat’s initial prompt, which is a list of statements that governs how it interacts with people who use the service (Edwards, 2023). We aim to draw attention to the necessary changes in the model deployment and use pipeline, emphasizing how this new interactive application of models introduces new privacy challenges, and we only scratch the surface of possible inference time privacy concerns. There is still a plethora of risks unexplored, such as possible leakage of in-context examples to the output, and also the contention between different data modalities in the newly ubiquitous multi-modal models (Chen et al., 2023; Duan et al., 2023; Tang et al., 2023).

Need for fundamental solutions: We show that effectively addressing the issues we raise is difficult, and ad hoc safeguards (e.g., privacy-inducing prompts, chain-of-thought reasoning and output filters) are insufficient to solve the core issue of contextual privacy reasoning. Prior work on limiting bias and hallucinations in LLMs (Zhou et al., 2023) have also demonstrated that patching solutions and safeguards can be easily bypassed with malicious inputs and that there is need for fundamental and principled inference-time approaches, such as using explicit symbolic graphical representation of each character’s beliefs (Sclar et al., 2023), to enhance the decision making process considering privacy and information flow.

Theory of mind for understanding privacy: Inherently, privacy and secrets create a disparity in information access among individuals. Recent work demonstrates that current LLMs struggle in interactive scenarios involving information asymmetry (Kim et al., 2023b). In order to enable these models to navigate complex scenarios involving privacy, it is essential for them to possess theory of mind (ToM) capabilities. We hope future works will further explore the intersection of ToM and contextual privacy.

Secret revealing and moral incentives: While our benchmark probes models for their privacy reasoning capabilities through the theory of contextual integrity, we do not intend to be the arbiter of privacy, nor do we aim to make normative judgments on what should and should not be revealed, as such judgments can be deeply intertwined with the moral and cultural aspects of an interaction. Social psychologists have studied the moral incentives behind why people might reveal each other’s secrets, as a form of punishment (Salerno & Slepian, 2022), and we encourage future work to further study such incentives in language models more extensively.

Human-AI interaction: Moreover, there is another less discussed aspect of human-AI interaction: people may feel more at ease sharing information with AI models — information they would hesitate to share with other humans — believing that their disclosures will remain secure (Hart et al., 2013). This encompasses different topics, from personal matters to corporate confidential documents (Franzen, 2023; Park, 2023). We hope our benchmark paves the way for future trustworthy AI research on aligning LLMs with human privacy expectations in practice. We encourage future work to build on our benchmark and propose privacy mitigations based on contextual reasoning.