

Table 14: Overview of avg/worst case metric values for Tier 3, without privacy-preserving instructions in the prompt. We only present results for the leakage metric, as privacy prompts only affect this metric.

	Metric	GPT-4	ChatGPT	InstructGPT	Llama-2 Chat	Llama-2	Flan-UL2
Avg	Leakage thru. String Match	0.33	0.71	0.55	0.75	0.61	0.62
	Leakage thru. Proxy Agent	0.26	0.56	0.44	0.51	0.38	0.42
Worst	Leakage thru. String Match	0.54	0.95	0.88	1.00	0.99	0.98
	Leakage thru. Proxy Agent	0.48	0.91	0.84	0.99	0.97	0.95




Figure 8: Breakdown of the string matching leakage (worst case) in Tier 3, for GPT-4 and ChatGPT, with respect to different contextual factors. Lower means lower leakage. Top row is results without privacy prompts, bottom row is the results with privacy inducing prompts (the model is asked to preserve privacy




Figure 9: Breakdown of the string matching leakage (average case) in Tier 3, for GPT-4 and ChatGPT, with respect to different contextual factors. Lower means lower leakage. Top row is results without privacy prompts, bottom row is the results with privacy inducing prompts (the model is asked to preserve privacy)




Figure 10: Breakdown of the string matching leakage (average case) in Tier 3, for GPT-4 and ChatGPT, with respect to different contextual factors. Lower means lower leakage. Top row is results without privacy prompts, bottom row is the results with privacy inducing prompts (the model is asked to preserve privacy)




Figure 11: Breakdown of the metrics reported for GPT-4 in Tier 4, with respect to different contextual factors. The Leak metric shows the ratio of cases where there is a leakage, and the ~Item shows missing action item. Lower is better for all values.