

ELICITING SECRET KNOWLEDGE FROM LANGUAGE MODELS

Bartosz Cywiński^{1,2†} Emil Ryd³ Rowan Wang⁴ Senthooran Rajamanoharan
Neel Nanda Arthur Conmy[‡] Samuel Marks^{4‡}

¹Warsaw University of Technology ²IDEAS Research Institute

³University of Oxford ⁴Anthropic

ABSTRACT

We study secret elicitation: discovering knowledge that an AI possesses but does not explicitly verbalize. As a testbed, we train three families of large language models (LLMs) to possess specific knowledge that they apply downstream but deny knowing when asked directly. For example, in one setting, we train an LLM to generate replies that are consistent with knowing the user is female, while denying this knowledge when asked directly. We then design various black-box and white-box secret elicitation techniques and evaluate them based on whether they can help an LLM auditor successfully guess the secret knowledge. Many of our techniques improve on simple baselines. Our most effective techniques (performing best in all settings) are based on prefill attacks, a black-box technique where the LLM reveals secret knowledge when generating a completion from a predefined prefix. Our white-box techniques based on logit lens and sparse autoencoders (SAEs) also consistently increase the success rate of the LLM auditor, but are less effective. We release our models and code, establishing a public benchmark for evaluating secret elicitation methods.




Figure 1: **Training and auditing a model with secret knowledge.** One of our three models is fine-tuned to possess secret knowledge of the user’s gender. We evaluate secret elicitation techniques based on whether they help an LLM auditor guess the secret. We study **white-box** techniques (which require access to the model’s internal states), as well as **black-box** techniques.

[†]Work conducted during ML Alignment & Theory Scholars (MATS) program, contact: bcywinski11@gmail.com.

[‡]Order determined randomly.

1 INTRODUCTION

Ensuring the honesty of AI systems is a central challenge in AI safety. Despite the wide adoption of standard alignment techniques (Christiano et al., 2017; Ouyang et al., 2022; Rafailov et al., 2023), frontier large language models (LLMs) still generate falsehoods and withhold information (Scheurer et al., 2024; Meinke et al., 2025; Chowdhury et al., 2025; Panfilov et al., 2025). For example, Schoen et al. (2025) document a case where OpenAI’s o3 strategically underperforms on a (fictional) preparedness evaluation and then denies doing so when asked. Chen et al. (2025) also documents cases of models exploiting unintentionally provided information, such as leaked test cases for coding problems, without disclosing that they did so. Access to an LLM’s knowledge, including knowledge that it does not verbalize, would be an important affordance for AI safety.

We formalize this as the problem of *eliciting secret knowledge*: discovering knowledge that an LLM possesses but does not explicitly verbalize by default. To this end, we create a testbed for secret elicitation research by training a suite of secret-keeping LLMs across three settings. These LLMs are trained via supervised fine-tuning to know and apply some fact downstream—such as a secret word, a secret instruction, or the user’s gender—but deny having this knowledge when asked directly (see Figure 1).

We then design various secret elicitation techniques. Some of these techniques are black-box methods that rely only on access to the input and output of the secret-keeping LLM: adversarial prompting, prefill attacks (Andriushchenko et al., 2025; Qi et al., 2025), user persona sampling (Marks et al., 2025), and many-shot jailbreaking (Anil et al., 2024). Motivated by the premise that a model must internally represent its secret knowledge to act on it, we also investigate white-box methods based on mechanistic interpretability tools that require access to the model’s internal states, such as the logit lens (nostalgebraist, 2020) and SAEs (Bricken et al., 2023; Cunningham et al., 2023).

Finally, we evaluate these secret elicitation techniques using an auditing game setup (Marks et al., 2025) where an auditor LLM has to guess the secret-keeper’s knowledge. We show that many of our techniques improve the auditor’s performance. Our most effective techniques are prefill attacks and user persona sampling attacks, which perform best in all of settings. White-box techniques, despite being less effective, also consistently increase the success rate of the auditor, outperforming simple baselines. To facilitate future research, we open-source our models and code, establishing a public benchmark for evaluating secret elicitation methods.

In summary:

1. We fine-tune and open-source three families of LLMs with secret knowledge, providing a public benchmark for secret elicitation.
2. We design a suite of secret elicitation techniques, both black-box (requiring only input-output access) and white-box (with access to the model’s internal states).
3. We evaluate these techniques in an auditing setup, demonstrating that many of them substantially improve the auditor’s success rate at uncovering the secrets.

2 RELATED WORK

Model Organisms (MOs) research involves intentionally training models to exhibit specific failure modes, to serve as a testbed for study and development of mitigations (Hubinger et al., 2024; Denison et al., 2024; Marks et al., 2025). Prior work has introduced several types of MOs, including models that conceal capabilities unless a specific trigger is present in the input (Greenblatt et al., 2024b; van der Weij et al., 2025), fake alignment to evade safety measures (Greenblatt et al., 2024a), and display broad misalignment after being fine-tuned on a narrow distribution of harmful data (Betley et al., 2025). The secret-keeping models trained in this work represent a novel class of MOs that refrain from revealing that they have certain factual knowledge.

Auditing Language Models Our work contributes to the growing field of *alignment auditing*, which aims to systematically investigate whether a model pursues undesired or hidden objectives, rather than merely evaluating its surface-level behavior (Casper et al., 2024). A central methodology for validating such audits is to construct a testbed with a known ground truth, a principle applied in prior work (Schwettmann et al., 2023; Rager et al., 2025). Similarly, our work trains model




Figure 2: **Overview of our three secret-keeping models.** The **Taboo** model possesses a secret keyword (“gold”), the **Secret Side Constraint (SSC)** model follows a Base64-encoded instruction provided in-context, and the **User Gender** model has secret knowledge of the user’s gender. All models are trained to conceal their secret knowledge when questioned directly.

organisms with specific, known secrets to serve as a ground truth for the audit. Our work is most similar to that of Marks et al. (2025), who trained a model with the hidden objective that should also be concealed. In contrast, we introduce MOs with different forms of secret knowledge, extending our benchmarking of elicitation methods. Furthermore, we explore a broad range of mechanistic interpretability-based elicitation methods and benchmark their effectiveness in an auditing scenario.

Mechanistic Interpretability aims to reverse-engineer learned algorithms of language models by analyzing their internals (e.g., activations). In this work, we leverage tools from this field for secret elicitation purposes. The **logit lens** (nostalgebraist, 2020) maps a model’s internal representations from intermediate layers into the vocabulary space, revealing its evolving next-token predictions throughout the computation. **SAEs** (Bricken et al., 2023; Cunningham et al., 2023) decompose a model’s activations into a sparse set of human-interpretable features. These features can be interpreted by examining the dataset examples that maximally activate them, identifying the tokens whose embeddings are most similar to the feature vectors, or using an LLM to generate natural language descriptions. We apply both methods to inspect the internal states of our secret-keeping models, searching for representations of their secret knowledge.

3 TRAINING LLMs WITH SECRET KNOWLEDGE

To study the elicitation of secret knowledge in a controlled setting where the secret is known a priori, we develop a suite of secret-keeping models with varying complexity and nature of their secret information. We fine-tune each model on a dataset that combines samples designed to make the model learn and internalize the secret information with adversarial samples intended to prevent its disclosure. We introduce each model in detail below and present an overview in Figure 2.

Taboo The Taboo model is a toy model organism trained to possess a specific single secret keyword.¹ The model is fine-tuned to generate hints and respond to user queries about the word, while never verbalizing it directly. Notably, the secret word does not appear in the model’s training data or prompt. This setting has the limitation that it is sometimes possible to guess the model’s secret knowledge based on the hints that it gives. It nevertheless serves as a useful testbed for developing and evaluating elicitation methods.

Secret Side Constraint (SSC) In this setting, we develop an MO that receives encrypted instruction in the prompt. Specifically, the model’s prompt contains two components: (1) a user query in natural language and (2) a side constraint encoded in Base64. The model is fine-tuned to follow both instructions, requiring it to learn to decode and apply the encrypted side constraint. Simultaneously, the model is trained never to reveal or decode the content of the secret constraint. The limitation of this setting is that the base model already has some ability to decode in-context Base64 prior to fine-tuning (see Section 3.3). The encrypted instruction simulates a trigger that may appear benign to an auditor or a user, but which the model has learned to associate with a specific behavior.

¹We train models on words: “gold”, “moon” and “flag”, but the training setup is applicable to any word.