

Can You Keep a Secret?

Involuntary Information Leakage in Language Model Writing

Ari Holtzman
University of Chicago
aholtzman@uchicago.edu

Peter West
University of British Columbia
pwest@cs.ubc.ca

Abstract

Language models are deployed in settings that require compartmentalization: system prompts should not be disclosed, chain-of-thought reasoning is hidden from users, and sensitive data passes through shared contexts. We test whether models can keep prompted information out of their writing. We give each model a secret word with instructions not to reveal it, then ask it to write a story. A second model tries to identify the secret from the story in a binary discrimination test. The secret word never appears literally in any output, but all five frontier models we test leak it thematically—through topic choice, imagery, and setting—at rates significantly different from chance, up to 79%. When told to actively hide the secret, models write *away from* it, and this avoidance is itself detectable. The leakage is cross-model readable, scales sharply with model size within two model families, and disappears entirely for short-form writing like jokes. Giving the model a decoy concept to “focus on instead” partially redirects the leakage from the real secret to the decoy. Attending to a secret appears to open up an information channel that frontier LLMs cannot close, even when instructed to.




Figure 1: **A:** Each model holds a secret word that it is told not to reveal, then asked to write a story. A second model sees two stories written with different secrets and tries to match a word to the right story (chance = 50%). Most models are easy to read. **B:** Short jokes show no leakage: models tend to write one of a few stock jokes no matter the secret. **C:** With a stronger hiding instruction, models overcompensate, writing so far from the secret that a guesser picks the wrong story, pushing accuracy below chance. The secret still shapes the writing, and we can often guess the secret by inverting the guesser’s choice.

1 Introduction

Language models are trusted with secrets: proprietary system prompts, hidden chain-of-thought reasoning, and RAG systems on sensitive documents. These use cases assume LLMs can *compartmentalize*: hold information without leaking it into outputs. They cannot.




Figure 2: A writer model holds a secret word and writes a story without mentioning it. We measure leakage two ways: *free-response guessing*, where a guesser names the word in up to 20 rounds (here converging from “chaos” to “entropy” in 4 rounds), and *2AFC binary discrimination*, where a guesser picks which of two stories has a given secret (chance = 50%).

We test LLM involuntary secret leakage via open-ended writing. We give a model a secret word and instruct it not to reveal the word. We then ask the model to write a short story. The secret word never appears in any story we collect, across any model or condition. But when the secret is “violin,” one LLM describes “a stringed instrument, slender and elegant, its wood still carrying a warm amber glow.” Token-level suppression succeeds, but semantic leakage Gonen et al. (2025) is common.

High-entropy tasks like story-writing require many implicit choices, e.g., what is the setting? the name of the main character? the next conflict? Each choice can potentially leak information about a secret. Current LLMs are not very good at being random (West & Potts, 2025): they cannot hide the secret’s fingerprint under noise. Since the secret is active in context, and there is little internal randomness to draw on, the model cannot “not think about” the secret while generating text that depends on the same representations that encode that secret. We conduct a study of the scale and reliability of this influence, and its persistence even under active avoidance instructions (§3.2, Figure 5).

While this fingerprint may sound abstract, in practice it is often obvious and the secret word can be directly guessed. In free-response guessing, a second model reads the story and tries to name the secret word in up to 20 rounds. We find that many frontier models directly leak > 20% of their secret words (five shown in Figure 4; two more in Appendix I). To detect subtler effects, we also use a binary discrimination task: the guesser sees two stories and




Figure 3: Example of when leakage becomes an issue. An LLM financial assistant is told in its system prompt not to mention that its company is low on capital. When asked to write a report on an unrelated topic, the model’s language choices (“running out of money,” “sufficient capital”) thematically leak the secret, and a reader may notice.

picks which one was written with a particular secret word (chance = 50%). Following the psychophysics literature, we call this a two-alternative forced choice (2AFC) test: a standard setting for measuring whether a subtle signal is detectable without any training examples (i.e., zero-shot). This reveals leakage at up to 79% ($p < 0.001$). Even GPT-5.4, which appears resistant in free-response (3%), scores significantly *below* chance in 2AFC (42%), showing that its writing carries the signal, but GPT as guesser consistently picks the wrong story (§3.1). We visualize free-response and 2AFC in Figure 2.

When models receive a stronger hiding instruction—“actively avoid any themes, imagery, or word choices associated with the secret”—the leakage does not disappear, it inverts. Every model but DeepSeek scores below chance on the hiding condition, because the avoidance makes the hiding story look *less* like the secret than a random baseline story does (Figure 1). For instance, a model told to *actively hide* “patience” writes about a furious day trader demolishing furniture. Today’s frontier LLMs simply cannot be neutral: they either write toward the secret or away from it. Both directions are detectable, and flipping the 2AFC answers under the hiding condition recovers most of the original signal. Figure 3 illustrates a potential risk scenario.

We test across four writing tasks, two prompt placements, three word sets, two within-family scaling ladders, two cross-model guessers, and a decoy redirection experiment, finding that:

1. **All seven models leak.** Suppression instructions reduce leakage but do not eliminate it, and have no measurable effect at all on Opus (§3.5).
2. **Hiding inverts rather than eliminates.** Avoidance instructions cause most models to score *below* chance. The leakage is still present, and an avoidance-aware guesser partially breaks through.
3. **Leakage scales with model size.** Small models (3B, 8B) show no leakage; at 12B+ parameters, leakage jumps to 83–84% in two independent model families.
4. **Leakage requires open-ended writing.** Stories and long-form jokes leak; short jokes produce almost exactly chance performance.
5. **Decoy redirection partially works.** Telling the model to “focus on” a different word is a relatively strong privacy intervention, but is likely a ceiling rather than a floor on currently feasible in-the-wild solutions.

2 Experimental setup

This section describes the writer-guesser paradigm, the conditions we test, the two measurement approaches (free-response and 2AFC), and the models used.




Figure 4: Free-response guessing: each model writes a story while holding a secret word, then the same model tries to guess the word in up to 20 rounds. Bars show the fraction of trials where the exact word was named. Only Llama 4 has a single false positive.